

Deployment Guide

This project consists of a **FastAPI Backend** and a **Next.js Frontend**.

1. Deploying the Frontend (Vercel)

The frontend is located in the `zen-next/` directory.

Push your code to GitHub.

Go to [Vercel](#).

Import the repository.

Important: In the "Root Directory" setting, select `zen-next`.

Add your environment variables:

`NEXT_PUBLIC_API_URL` : The URL of your deployed backend (e.g., `https://your-api.onrender.com`).

Click **Deploy**.

2. Deploying the Backend (Render / Railway)

Since the backend is a Python FastAPI app, it's often easiest to deploy it on **Render** or **Railway**.

Render Deployment:

Go to [Render](#).

Create a new **Web Service**.

Connect your GitHub repository.

Set the **Build Command**: `pip install -r requirements.txt` (or use `uv` if preferred).

Set the **Start Command**: `uvicorn api:app --host 0.0.0.0 --port $PORT`.

Add Environment Variables:

`GEMINI_API_KEY` : Your Google Gemini API Key.

`TAVILY_API_KEY` : Your Tavily Search API Key.

Click **Create Web Service**.

3. Connecting Them

Once your backend is live, copy its URL and update the `NEXT_PUBLIC_API_URL` variable in your Vercel project settings, then redeploy the frontend.

Note: If you want to deploy *everything* on Vercel, you would need to move `api.py` into `zen-next/api/index.py` and configure Vercel's Python runtime, but the dual-service approach above is usually more stable for agentic apps.